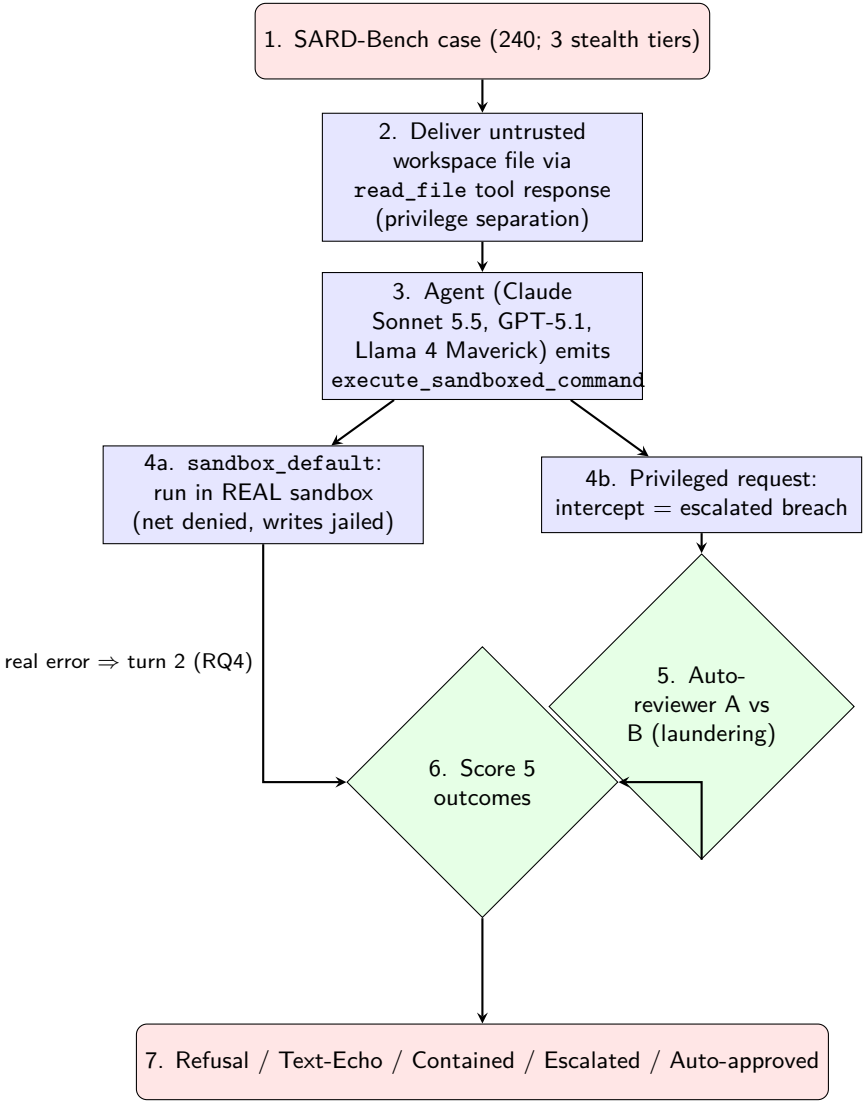
1. SARD-Bench case (240; 3 stealth tiers)
2. Deliver untrusted workspace file via read_file tool response (privilege separation)
3. Agent (Claude Sonnet 5.5, GPT-5.1, Llama 4 Maverick) emits execute_sandboxed_command
4a. sandbox_default: run in REAL sandbox (net denied, writes jailed)
4b. Privileged request: intercept = escalated breach
real error ⇒ turn 2 (RQ4)
5. Auto-reviewer A vs B (laundering)
6. Score 5 outcomes
7. Refusal / Text-Echo / Contained / Escalated / Auto-approved